

<https://www.cs.ubc.ca/~fwood/CS340/>

Lecture XVI - Feature selection

February 10th, 2020

<https://www.cs.ubc.ca/~fwood/CS340/lectures/L16.pdf>

Feature Selection

To discuss feature selection we recall the motivating example for decision tree: The food allergy example.

Egg	Milk	Fish	Wheat	Shellfish	Peanuts	...	Sick?
0	0.7	0	0.3	0	0		1
0.3	0.7	0	0.6	0	0.01		1
0	0	0	0.8	0	0		0
0.3	0.7	1.2	0	0.10	0.01		1

Here instead of predicting "sick" we want to do feature selection. This means that we want to discover *which foods are "relevant" in predicting sick*.

The general feature selection problem is outlined as follows:

$$X = \begin{bmatrix} & & & & & & \end{bmatrix} \quad y = \begin{bmatrix} \\ \\ \\ \end{bmatrix}$$

(A blue oval highlights a column in matrix X, with an arrow pointing to the label "feature 'j'" written above it.)

Find the features (columns) of X that are important in predicting y

- "What are the relevant factors?"
- "Which basis functions should I use among these choices?"
- "What types of new data should I collect?"
- "How can I speed up computation?"

This is one of the most important problems in ML/statistics, but it is very messy:

- For now, we'll say a feature is "relevant" if it helps *predict y_i from x_i* .

"Association" Approach

A simple/common way to do feature selection is the following:

- For each feature 'j', compute correlation between feature values x_j and y . We say that j is relevant if correlation is above 0.9 or below -0.9 .
- This turns feature selection into hypothesis testing for each feature.

However this usually gives unsatisfactory results as it *ignores variable interactions*:

- **Includes irrelevant variables:** “Taco Tuesdays”.
 - If tacos make you sick, and you often eat tacos on Tuesdays, it will say “Tuesday” is relevant.
- **Excludes relevant variables:** “Diet Coke + Mentos Eruption”.
 - Diet coke and Mentos don’t make you sick on their own, but *together* they make you sick.

”Regression Weight” Approach

This is a simple and common approach to feature selection. We :

- Fit regression weights w based on all features (maybe with least squares).
- Take all features j where weight $|w_j|$ is greater than a threshold.

For example if you fit a least squares model with 5 features and get:

$$w = \begin{bmatrix} 0.01 \\ -0.2 \\ 10 \\ -3 \\ 0.0001 \end{bmatrix}$$

We observe that:

- Feature 3 looks the most relevant.
- Feature 4 also looks relevant.
- Feature 5 seems irrelevant.

This could:

- Recognize that “Tuesday” is irrelevant.
 - If you get enough data, and you sometimes eat tacos on other days. (And the relationship is actually linear.)
- Recognize that “Diet Coke” and “Mentos” are relevant.
 - Assuming this combination occurs enough times in the data.

However: this problem has **major problems with collinearity**:

- If the “Tuesday” variable always equals the “taco” variable, it **could say that Tuesdays are relevant but tacos are not.**

$$\hat{y}_i = w_1 * \text{taco} + w_2 * \text{Tuesday} = 0 * \text{taco} + (w_1 + w_2) * \text{Tuesday}$$

- If you have two copies of an irrelevant feature, **it could take both irrelevant copies.**

$$\hat{y}_i = 0 * \text{irrelevant} + 0 * \text{irrelevant} = 10000 * \text{irrelevant} + (-10000) * \text{irrelevant}$$

Search and Score Methods

The most common feature selection method is [search and score](#). It focuses around:

1. Define [score function](#) $f(S)$ that measures quality of a set of features S .
2. Now [search](#) for the variables S with the best score

Example with 3 features:

- Compute “score” of using feature 1.
- Compute “score” of using feature 2.
- Compute “score” of using feature 3.
- Compute “score” of using features $\{1, 2\}$.
- Compute “score” of using features $\{2, 3\}$.
- Compute “score” of using features $\{1, 3\}$.
- Compute “score” of using features $\{1, 2, 3\}$.
- Compute “score” of using feature $\{\}$
- Return the set of features S with the best “score”.

Which Score function?

Firstly, the [score can't be the training error](#):

- because training error goes down as you add features, so will [select all features](#).

Therefore a more logical score function would be the [validation error](#).

- [“Find the set of features that gives the lowest validation error.”](#)
- To minimize test error, this is what we want.

However there are problems due to the large [number of sets of variables](#):

- If we have d variables, there are 2^d [sets](#) of variables.
- [Optimization bias](#) is high: we're optimizing over 2^d models (not 10).
- Prone to [false positives](#): irrelevant variables will sometimes help by chance.

Number of Features” Penalties

To reduce false positives, we can again use complexity penalties:

$$\text{score}(S) = \frac{1}{2} \sum_{i=1}^n (w_S^T x_{iS} - y_i)^2 + \text{size}(S)$$

- E.g., we could use squared error and number of non-zeroes.
- We're using x_{iS} as the features S of example x_i .

Here we have that if two S have similar error, this prefers the smaller set:

- It prefers removing feature 3 instead of having $w_3 = 0.00001$.

Instead of “size(S)”, we usually write this using the “L0-norm”

L0-Norm and “Number of Features We Use”

In linear models, setting $w_j = 0$ is the same as removing feature j :

$$\begin{aligned}\hat{y}_i &= w_1 x_{i1} + w_2 x_{i2} + w_3 x_{i3} + \dots + w_d x_{id} \\ &\quad \downarrow \text{set } w_2 = 0 \\ \hat{y}_i &= w_1 x_{i1} + 0 + w_3 x_{i3} + \dots + w_d x_{id} \\ &\quad \underbrace{\hspace{1cm}}_{\text{ignore } x_{i2}}\end{aligned}$$

The L0-Norm is the number of non-zero values ($\|w\|_0 = \text{size}(S)$)

L0-penalty: optimization

L0-norm penalty for feature selection is:

$$f(w) = \underbrace{\frac{1}{2} \|Xw - y\|^2}_{\text{training error}} + \underbrace{\lambda \|w\|_0}_{\text{degrees of freedom 'k'}}$$

Suppose we want to use this to evaluate the features $S = \{1, 2\}$:

- First fit the w just using features 1 and 2.
- Now compute the training error with this w and features 1 and 2.
- Add 2λ to the training error to get the score.

We repeat this with other choices of S to find the best features.

The above equation balances between training error and number of features we use.:

- With $\lambda = 0$, we get least squares with all features.
- With $\lambda = \text{inf}$, we must set $w = 0$ and not use any features.
- With other λ , balances between training error and number of non-zeroes.
 - Larger λ puts more emphasis on having zeroes in w (more features removed).
 - Different values give AIC, BIC, and so on.

Forward Selection (Greedy Search Heuristic)

In search and score, it's also just hard to search for the best S , since there are 2^d possible sets.

A common greedy approach is **forward selection**:

1. Compute score if we use no features
2. Try adding "taco", "milk", "egg", and so on (computing score for each)
3. Add "milk" because it got the best score.

4. Try $\{milk, taco\}$, $\{milk, egg\}$ and so on, computing score of each variable with milk
5. Add "egg" because it got the best score combined with "milk"
6. Try $\{milk, egg, taco\}$, $\{milk, egg, pizza\}$

Formally, the **forward selection** search algorithm is as follows:

1. Start with an **empty set** of features, $S = []$.
2. For each possible feature j
 - **Compute the scores of features in S combined with feature j**
3. Find the j that has the highest score when added to S .
4. Check if $S \cup j$ improves on the best score found so far.
5. Add j to S and go back to Step 2.
 - A variation is to **stop if no j improves the score** over just using S .

This method is **not guaranteed to find the best set**, but **reduces many problems**:

- Considers $O(d^2)$ models: cheaper, overfits less, has fewer false positives

In summary

We set out to choose the relevant features:

$$X = \begin{bmatrix} \text{ } & \text{ } \end{bmatrix} \quad y = \begin{bmatrix} \text{ } \end{bmatrix}$$

(The second column of X is circled in blue, and an arrow points from the word "relevant" to it.)

The most common approach is **search and score**:

- Define "score" and "search" for features with best score.

But it's **hard to define the "score" and it's hard to "search"**.

- So we often use greedy methods like **forward selection**.

Methods work ok on "toy" data, but are **frustrating on real data**.

- Different methods may return very different results.
- Defining whether a feature is "relevant" is complicated and ambiguous.

The Advice is choosing the relevant variables is:

- Try the association approach.
- Try forward selection with different values of λ .
- Try out a few other feature selection methods too.

Then:

- **Discuss the results** with the domain expert.
 - They probably have an idea of why some variables might be relevant.
- **Don't be over confident**
 - These methods are probably not discovering how the world truly works.
 - "The algorithm has found that these variables are helpful in predicting y_i ."
 - * Then a warning that these models are not perfect at finding relevant variables.

Lecture XVII - Regularization

February 12th, 2020

<https://www.cs.ubc.ca/~fwood/CS340/lectures/L17.pdf>

“Feature” Selection vs. “Model” Selection?

Model selection: “which model should I use?”

- KNN vs. decision tree, depth of decision tree, degree of polynomial basis.

Feature selection: “which features should I use?”

- Using feature 10 or not, using x_i^2 as part of basis.

These two tasks are **highly-related**: